

2.4 · Data-Flow and Minimization Workshop

Week 2 · Tuesday Jul 21 — Data, Environment, and Human-Centered Sustainable AI
Generative AI & Sustainability · Munich Summer School 2026

Session goal

Ground the map in a real room, then shrink it: audit this classroom as your specimen space, state each data source's purpose, cross out what you won't collect, and commit to where inference and storage happen.

Step 0 · Audit the room you're sitting in (20 min)

Walk this classroom as if it were your learning space. Three columns on paper:

Already sensed	Could be sensed	We refuse to sense
what the building does now (HVAC? wifi APs counting devices?)	occupancy, CO ₂ , light, noise, plug loads — a €15 plug meter covers most of it	cameras? microphones? identity? seat-level tracking?

- Every "could" needs a **purpose** — no purpose, no sensor
- The "refuse" column is a pitch asset: it's your privacy story in one table

Steps (after the audit)

- Complete the inventory: add the non-sensor sources across your three agents — prompts and conversations, learning records, and external retrieval
- Apply data minimization: for each source state the purpose, and decide what should not be collected at all
- Make explicit private-ML choices: on-device vs cloud inference, what is logged and for how long, whether personal data enters training, and where DP, federation, or anonymization apply
- Keep sensors to the small, low-risk set that survived your audit

Deliverables by end of session

- Data-flow and minimization map (in progress), seeded by the room audit's three columns

Time and checkpoints

- 1:15–1:35 — **room audit**: walk the classroom; fill the three columns (already / could / refuse)
- 1:35–1:55 — inventory: add prompts, records, and retrieval across your three agents
- **1:55 checkpoint** — instructor pass: is anything missing? (chat logs and tool outputs are the usual blind spots)
- 1:55–2:10 — minimize: purpose per source, cross-outs, on-device vs. cloud, logging and retention decisions
- **2:10 share-out** — each team names one thing in its *refuse* column and why
- 2:10–2:15 — clean up the map for use in the next block